

# MARL-Based Trust Evaluation for Zero-Trust Low-Altitude UAV Networks

Anonymous Authors

Department of Computer Science and Engineering, Anonymous Institution

Correspondence: to be provided after peer-review anonymization is removed.

**Abstract:** Low-altitude UAV networks are characterized by rapidly changing topology, intermittent connectivity, and limited onboard computation and energy, which make continuous security assessment difficult in infrastructure-scarce missions. Under these conditions, compromised UAVs can conduct spoofing, data tampering, selective forwarding, dormant/on-off misbehavior, and limited collusion, and a high malicious-node ratio can distort local observations enough to make trust evaluation inaccurate. Existing evidence-fusion, anomaly-profile, and group-decision trust mechanisms can detect some insider behaviors, but their manually specified fusion logic often loses discrimination under dense or temporally evasive attacks; strong learned risk classifiers further narrow the detection gap but do not convert trust estimates into reward and penalty decisions that guide future UAV behavior. To address this problem, this paper proposes the Multi-Agent Trust and Incentive Mechanism (MTIM), a closed-loop framework that combines trust evaluation with a MARL-supported incentive mechanism for continuous zero-trust verification in low-altitude UAV networks. MTIM treats trust as a behavioral control signal: current UAV behavior is first evaluated through behavioral, communication, energy, task, and historical evidence; the resulting trust state is then used to assign rewards, penalties, and task-allocation preferences; and an offline-trained parameterized MARL policy strengthens this trust-to-incentive mapping under partial observations. Multi-agent execution allows different Thinker Agents to supervise local UAV subsets, while the Intelligent Decision Agent aggregates their estimates with confidence-aware uncertainty control. In the 50-seed trust-stream evaluation, MTIM achieves 94.4% detection accuracy, 91.8% F-score, and 90.5% task success rate under a 30% malicious-node ratio. Compared with paper-level reimplementations of MFA, FUBA, GALTrust, and GDM-DTM, MTIM improves both detection and task success under high malicious-node ratios; compared with a strong MLP-Trust baseline, MTIM obtains statistically similar detection performance but improves task success by 2.2 percentage points through incentive coupling. These results support the central claim that trust assessment is most effective when coupled to incentive-aware mission decisions, while packet-level multi-seed validation and real-flight deployment remain future work.

**Keywords:** UAV networks; trust management; multi-agent reinforcement learning; zero trust architecture; FANET security

## 1 Introduction

Low-altitude UAVs are becoming critical infrastructure for logistics, inspection, sensing, and emergency response, and their flexibility and rapid deployment make them a cornerstone of the emerging low-altitude economy [1]. The operating background of such networks is distinctive: flying ad hoc networks (FANETs) have rapidly changing topology, intermittent air-to-air connectivity, and severely constrained onboard computation and energy resources [2].

The security problem addressed in this paper is not merely that UAVs may be attacked, but that compromised nodes can remain inside the cooperative swarm and distort trust evidence over time. In practical missions, malicious UAVs may spoof state information, selectively forward packets, tamper with sensing or task results, manipulate energy/task behavior, or submit biased peer observations. When the malicious-node ratio is high, these behaviors can shift the observed trust distribution and make accurate trust assessment difficult, especially under dormant or on-off attacks that deliberately exploit historical trust.

To address this problem, this paper proposes the Multi-Agent Trust and Incentive Mechanism (MTIM). MTIM is introduced here as a behavior-guidance framework rather than only as a detector: it evaluates the current trustworthiness of each UAV, uses the trust score as the basis for rewards and

penalties, and feeds the resulting incentive decision back into task allocation so that cooperative behavior is favored over time.

Trust management has emerged as a core security primitive for UAV networks, enabling continuous assessment of node trustworthiness to detect malicious activities and mitigate attack impacts. Existing mechanisms have made meaningful progress, but the specific problem above remains only partially solved. FUBA uses fuzzy behavior analytics over heterogeneous FANET indicators, GALTrust introduces generative-adversarial/profile-based trust modeling, and GDM-DTM applies dynamic group decision making for low-altitude UAV trust evaluation [8, 10, 23]. These methods can detect certain insider or anomaly behaviors, so the unresolved gap is more specific: under dense malicious populations and temporal evasion, manually specified fusion rules, profile distances, or group-decision consistency penalties may lose discrimination, and their trust scores are not directly coupled to mission-level reward, penalty, or task-allocation decisions. Strong learned classifiers can further improve risk discrimination, but risk-only classifiers still leave the trust-to-action loop open.

This framing leads to three challenges. First, the trust evaluator must remain accurate when the malicious population is dense enough to distort local trust distributions. Second, the selected trust factors must be broad enough to capture coordinated deception, yet compact enough to avoid unnecessary computation. Third, the system must turn trust scores into concrete incentives: high-trust UAVs should receive favorable task or resource opportunities, whereas low-trust UAVs should be penalized, restricted, or excluded from sensitive task assignments. Zero trust offers the security principle for continuous verification [15], while recent UAV studies show that MARL is useful for distributed decision making under partial observability and dynamic coupling [16, 26–28]. In MTIM, however, reinforcement learning is used primarily to strengthen the incentive and trust-to-action policy, not to claim that raw trust evaluation requires reinforcement learning by itself.

**Problem statement.** This paper addresses the trust-layer decision problem in fixed-identity low-altitude UAV swarms: under dynamic topology, intermittent observations, and high malicious-node ratios, how can a trust mechanism accurately infer node risk from noisy multi-factor evidence and then convert the inferred trust state into mission-level incentive decisions without relying on centralized infrastructure or online labels? This scope deliberately focuses on continuous trust verification and task-level response, rather than complete identity management, Sybil-resistant admission control, or real-flight deployment.

As shown in Figure 1, MTIM implements the above idea as a distributed trust–incentive loop. Trust evaluation measures current behavior and computes a credibility score; reward and penalty actions use that score as their decision basis; the MARL-style policy improves how local evidence is translated into incentive choices under partial observation; and multiple Thinker Agents divide local supervision responsibilities before the Intelligent Decision Agent aggregates their estimates. The main contributions of this paper are as follows:

- **A trust-guided behavior-regulation framework:** MTIM constructs a closed loop in which UAV behavior is observed, converted into a trust score, and then used as the criterion for reward, penalty, and task-allocation decisions. This frames trust management as a mechanism for encouraging cooperative behavior, not only as a post hoc malicious-node detector.
- **A justified multi-factor trust assessment model:** We design a compact five-factor trust evaluator covering behavioral consistency, communication reliability, energy behavior, task execution, and historical reputation. These factors are selected because they correspond to distinct attack surfaces and temporal evidence sources; the model avoids adding unrestricted evidence dimensions, and the ablation study quantifies which factors contribute most to detection and task success.
- **A MARL-supported multi-agent incentive policy:** We use an offline-trained MARL-style policy to strengthen the mapping from trust states to reward/penalty actions, while multiple Thinker Agents provide local supervision under partial observability. The Intelligent Decision Agent then fuses local estimates with uncertainty-aware confidence weights. The empirical advantage is therefore mission-level: MTIM remains close to the strong MLP-Trust detector in accuracy but improves task success through incentive coupling.

The remainder of this paper is organized as follows. Section 2 reviews related work on trust management in UAV networks and MARL for network security. Section 3 presents the trust–incentive problem definition, network and threat models, multi-factor trust evaluation, MARL-supported incentive policy, and trust state dynamics. Section 4 describes the experimental setup and presents the performance

evaluation results. Section 5 concludes the paper and outlines future research directions.

## 2 Related Work

Trust evaluation for UAV networks now includes reputation models, evidence fusion, blockchain support, and learning-based designs. Yet low-altitude UAV studies rarely integrate zero-trust verification, distributed trust learning, uncertainty aggregation, and mission incentives. We review three related areas: UAV trust management, zero-trust assessment, and MARL for UAV networking and security.

### 2.1 Trust Management in UAV and FANET Systems

Recent surveys show that trust management has become a core security primitive in FANETs, where dynamic topology, intermittent links, and resource limits make purely identity-based protection insufficient [1–5]. Existing schemes fall into three categories: behavior-centric trust analytics, distributed or blockchain-supported trust, and learning-enhanced trust evaluation.

Behavior-centric trust evaluation widely uses fuzzy and evidence-fusion methods to combine heterogeneous observations. Benfriha et al. proposed FUBA, which uses fuzzy logic to integrate indicators such as RSSI, packet delivery ratio, delay, battery level, and humidity to identify insider threats in FANETs [8]. Li et al. developed an adaptive multi-granularity trust management scheme for UAV visual sensor security under adversarial attacks, where coarse- and fine-grained evidence is fused to improve robustness against known and unknown threats [9]. Akram et al. introduced GALTrust, which combines generative adversarial learning with type-2 fuzzy logic to improve adaptability to emerging malicious behaviors in spatial crowdsourcing drone services [10]. More recent federated-security studies also move trust assessment toward collaborative anomaly detection, including FL-based intrusion detection for FANETs, multimodal swarm anomaly detection, and communication-efficient FL for resource-constrained UAV-IoT systems [11–13]. These studies significantly enrich trust evidence modeling. Their remaining limitation for the problem considered here is not that they cannot detect insider behavior at all, but that fuzzy, profile-based, or task-specific fusion can become sensitive to distribution shifts when the malicious population is dense and attack timing is deliberately evasive.

A second line of work improves trust persistence and auditability through distributed storage or hierarchical trust exchange. Ge et al. developed a provenance-aware distributed trust model for resilient UAV networks, highlighting the value of peer-to-peer evidence tracing in hostile environments [7]. Zuo et al. proposed a hierarchical blockchain-based trust measurement method for drone clusters to address trust inheritance, cross-cluster re-evaluation, and historical trust verification under dynamic node entry and exit [6]. Akram et al. proposed TMIoDT, which integrates digital twins, blockchain, and deep-learning-based intrusion detection to secure Open RAN-based Internet of Drone Things applications [17]. Closely related secure-FL frameworks have further explored malicious-node detection and trusted collaboration through blockchain-supported service provisioning, privacy-preserving IoDT security, and hierarchical federated aggregation [18–21]. Related work on secure quantum federated learning in UAV-assisted wireless networks further indicates that trust-enhanced game incentives can improve participation quality and strengthen learning robustness under strategic behavior [22]. More recently, Hu et al. proposed GDM-DTM for malicious-node detection in low-altitude UAV networks by combining group decision making with dynamic trust management [23]. These schemes improve distributed credibility, trust persistence, and malicious-node detection. For our setting, the central gap is that the trust state is usually treated as an assessment output rather than as the basis for learned reward, penalty, and task-allocation decisions.

### 2.2 Zero-Trust and Continuous Trust Assessment

Zero trust shifts security from perimeter-based assumptions to continuous, resource-centric verification. The NIST zero-trust architecture defines this paradigm through the principle of “never trust, always verify” and replaces implicit perimeter-based trust with continuous, context-aware verification [15]. Building on this idea, Wang et al. proposed a zero-trust-driven collaborative intrusion detection framework for IoT that uses continuous trust assessment to support distributed security decisions [24]. In the UAV and 6G context, Al Ridhawi and Aloqaily integrated zero trust, digital twins, federated learning, and blockchain into a UAV-enabled 6G framework, highlighting the broader relevance of decentralized trust enforcement in highly dynamic aerial networks [25]. In the UAV authentication domain, Wang et al.

introduced a context-aware trust evaluation model for malicious UAV detection during authentication, showing that trust can also be injected into the authentication stage rather than being treated purely as a post hoc reputation signal [14].

Zero trust is particularly relevant to low-altitude UAV swarms: node roles may change during flight, links may appear and disappear rapidly, and local observations may be incomplete or noisy. Trust decisions cannot be based solely on prior membership or static credentials, but must be continuously updated from current observations, device status, communication behavior, and mission context. Therefore, trust decisions should be dynamic, uncertainty-aware, and resource-conscious rather than static or binary.

These studies move trust evaluation closer to continuous verification. Nevertheless, the current zero-trust literature still focuses mainly on authentication, access control, or intrusion detection. It does not fully address how trust should evolve online in low-altitude UAV swarms when heterogeneous evidence from behavior, communication quality, task execution, energy state, and historical interactions must be fused under strict resource constraints. In particular, uncertainty-aware aggregation and decentralized incentive adaptation remain underexplored in zero-trust UAV settings.

## 2.3 MARL for UAV Networking and Security

MARL has recently become a prominent framework for decentralized decision making in UAV systems. A recent survey by Ekechi et al. reports that MARL has mainly been used for UAV control, formation, coverage, coordination, and cooperative optimization, rather than for trust management itself [16]. Representative examples include the MADDPG-based multi-UAV trajectory planning framework of Wang et al. for UAV-assisted mobile edge computing [26], the semi-distributed multi-agent federated reinforcement learning approach of Nie et al. for resource management in UAV-aided MEC systems [27], and the joint optimization of trajectory control, task offloading, and resource allocation in air-ground integrated networks by Alam and Moh [28]. Closely related learning-and-optimization studies in UAV-enabled federated settings have examined asynchronous scheduling, contract-based participant selection, hierarchical personalized learning, air-ground FL coordination, and energy-aware trajectory and resource management [29–34]. Together, these cited works provide evidence that distributed learning can support optimization under partial observability and dynamic coupling among UAVs. More recently, Cao et al. explored LLM-based multi-agent situation awareness for zero-trust space-air-ground integrated networks [35]. Their results further suggest that model selection in security-oriented multi-agent systems should balance capability and deployment cost; for specialized subtasks, smaller task-adapted models can be more practical than larger general-purpose models.

However, existing MARL applications in UAV networking remain largely performance driven, with objectives such as delay reduction, energy efficiency, coverage, or offloading efficiency. The reward design in these studies is rarely aligned with trustworthiness, collusion resistance, or long-horizon security incentives. MARL has been used to optimize how UAVs move, communicate, or allocate resources, but not how they collaboratively evaluate trust, calibrate trust-factor importance, and propagate trustworthy decisions under zero-trust assumptions.

Overall, prior work has made strong progress in trust modeling, distributed auditability, and UAV-oriented MARL, but three problem-specific challenges remain insufficiently addressed for low-altitude UAV trust evaluation:

- (1) **Accurate trust assessment under high malicious-node ratios:** Fuzzy, profile-based, and group-decision mechanisms can lose discrimination when the malicious population substantially shifts local trust distributions, leading to poor mission continuity.
- (2) **Temporal and multi-factor evasion:** Dormant and on-off attackers can exploit historical trust accumulation, while multi-factor manipulation requires deciding which behavioral, communication, energy, task, and history evidence remains reliable at each update.
- (3) **Closing the trust-to-action loop:** A detector that outputs only a trust or risk score does not by itself decide how task allocation, incentives, or penalties should change; this is the specific gap exposed by the comparison with the strong risk-only MLP-Trust baseline.

MTIM addresses these three challenges by coupling multi-factor trust evaluation with a learned reward–penalty policy, multi-agent local supervision, confidence-aware aggregation, and temporal trust decay.

### 3 System Model

We consider a low-altitude UAV swarm executing a cooperative mission (e.g., area surveillance, delivery, or search-and-rescue) in an open, infrastructure-scarce environment. The network is subject to three intrinsic constraints—dynamic topology, intermittent connectivity, and limited onboard resources—that simultaneously degrade conventional perimeter-based defenses and amplify the impact of insider threats. When malicious nodes constitute a non-negligible fraction of the swarm, sparse observations make accurate trust inference a problem of *reasoning under uncertainty* rather than deterministic rule-checking. This section first defines the trust-incentive decision problem (Section 3.1), then formalizes the network model (Section 3.2), threat model (Section 3.3), multi-factor trust evaluation mechanism (Section 3.4), MARL-style incentive policy (Section 3.5), and trust state dynamics (Section 3.5.4).

#### 3.1 Problem Definition


At each update interval, the system receives partial and noisy observations of each UAV’s behavior, communication quality, energy state, task execution, and historical reputation. The objective is to infer a trust state that reflects the UAV’s current behavior and then use that state to regulate future participation through rewards, penalties, and task-allocation preferences. Formally, for each UAV  $U_k$ , MTIM must output a global trust score  $\mathcal{T}_k(t)$ , a trust class  $C_k(t)$ , and an incentive action  $\pi_{jk}(t)$  for the local Thinker Agent  $j$  that supervises  $U_k$ :

$$(\mathcal{T}_k(t), C_k(t), \pi_{jk}(t)) = \Pi_\theta(o_j(t), \mathcal{H}_k(t)), \quad (1)$$

where  $o_j(t)$  is the partial observation available to agent  $j$ ,  $\mathcal{H}_k(t)$  is the recent trust-evidence history of node  $k$ , and  $\Pi_\theta$  is the learned trust-incentive policy. The conceptual objective is to improve both malicious-node detection and mission continuity while respecting update-cycle resource budgets:

$$\max_{\Pi_\theta} \mathbb{E}[\lambda_1 F_1 + \lambda_2 \text{TSR} - \lambda_3 C_{\text{comm}} - \lambda_4 T_{\text{eval}}], \quad (2)$$

subject to bounded trust scores  $\mathcal{T}_k(t) \in [0, 10]$  and the incentive budget later defined in Eq. 22. In the empirical implementation, this objective is approximated by a supervised surrogate loss (Section 3.5.3) rather than optimized directly. This distinction is important: the trust factors can be computed without reinforcement learning, while reinforcement learning is used to strengthen the mapping from trust evidence to reward-penalty actions under partial observation.



**Figure 1:** Overview of the MTIM framework. Summarization Agents collect UAV-network and ground-control observations, verified trust branches T1–T5 extract multi-factor evidence, parallel Thinker Agents evaluate locally observed UAV subsets and propose incentive actions, and the Intelligent Decision Agent aggregates trust estimates and closes the reward-penalty feedback loop.

### 3.2 UAV Network Model

We model the low-altitude UAV network as a dynamic ad hoc network of  $N$  heterogeneous UAV nodes,  $\mathcal{U} = \{U_1, U_2, \dots, U_N\}$ . Each node  $U_i$  is characterized by a tuple  $\langle \mathbf{p}_i(t), v_i(t), E_i(t), C_i \rangle$ , where  $\mathbf{p}_i(t) \in \mathbb{R}^2$  is the 2D position at time  $t$ ,  $v_i(t) \in [v_{\min}, v_{\max}]$  is the velocity,  $E_i(t) \in [0, E_i^{\max}]$  is the residual energy, and  $C_i$  denotes the computational capacity (CPU cycles/s). Nodes may enter or leave the swarm during the mission according to operational requirements.

The communication topology is a time-varying undirected graph  $G(t) = (\mathcal{U}, \mathcal{E}(t))$ , where an edge  $(i, j) \in \mathcal{E}(t)$  exists iff the Euclidean distance satisfies  $d_{ij}(t) = \|\mathbf{p}_i(t) - \mathbf{p}_j(t)\|_2 \leq R_{\text{comm}}$ , with  $R_{\text{comm}} = 300$  m (Table 2). Each link  $(i, j)$  has a time-varying quality characterized by packet delivery ratio  $\text{PDR}_{ij}(t)$ , received signal strength  $\text{RSSI}_{ij}(t)$ , and one-way latency  $L_{ij}(t)$ . The network is *infrastructure-scarce*: no fixed base stations or ground controllers are assumed to be continuously reachable; all trust computation and decision-making must execute on-board or via peer-to-peer exchange.

### 3.3 Threat Model

We adopt a white-box adversary model that is standard in the trust management literature [8, 23] but extend it to capture coordinated multi-dimensional manipulation.

**Adversary capabilities.** A subset of nodes  $\mathcal{M} \subset \mathcal{U}$  ( $|\mathcal{M}|/N \in [0.05, 0.50]$  in our experiments) is controlled by an adversary. Each malicious node  $U_m \in \mathcal{M}$  can independently manipulate any subset of the five trust dimensions by:

- **T1 (Behavioral)**: deviating from planned flight paths, skipping scheduled heartbeats, or responding inconsistently to coordination messages;
- **T2 (Communication)**: performing selective forwarding (dropping a fraction  $p_{\text{drop}} \in [0.3, 1.0]$  of packets), injecting spurious traffic, or manipulating transmission power to create asymmetric links;
- **T3 (Energy)**: exhibiting anomalous consumption patterns (e.g., sudden depletion through jamming, or unusually low consumption by remaining idle during active mission phases);
- **T4 (Task)**: returning fabricated or degraded task results (e.g., falsified sensor readings, partial payload delivery);
- **T5 (Historical)**: this dimension is not directly manipulable by the adversary—it is an accumulated reputation signal maintained by MTIM—but its update is indirectly affected when the adversary manipulates T1–T4.

**Attack strategies.** We consider three canonical attack patterns that cover the spectrum from naive to evasive adversaries:

- (i) **Constant attack**: the malicious node exhibits adversarial behavior from  $t = 0$  and maintains it throughout the mission. This models unsophisticated attackers and serves as a lower bound on detection difficulty.
- (ii) **Dormant attack**: the malicious node behaves honestly during an initial warm-up period ( $t \leq T_{\text{warm}}$ ), accumulating trust, then switches to adversarial behavior. This models adversaries that attempt to evade initial screening.
- (iii) **On-off attack**: the malicious node toggles between honest and adversarial behavior with period  $T_{\text{toggle}}$ , attempting to keep its trust score near the classification boundary where decisions are most uncertain.

In mixed-attack experiments, each pattern accounts for one-third of the malicious population. The main evaluation focuses on direct behavioral manipulation in T1–T4. To clarify the boundary of the zero-trust claim, we also include a limited collusion stress test in the sensitivity analysis: a subset of malicious nodes submits inflated peer recommendations while continuing one of the three behavioral attack patterns. Full Sybil resistance, where a physical adversary creates many unlinkable identities, requires identity-binding and admission-control mechanisms outside the scope of MTIM and is treated as a deployment prerequisite rather than a solved problem.

**Adversary knowledge.** The adversary is assumed to know the set of trust dimensions used by MTIM but *not* the frozen policy parameters, per-node policy outputs  $\omega_m(t)$ , or the Intelligent Decision Agent confidence weights  $\phi_{jk}(t)$ . This models a capable but not omniscient attacker. The adversary may submit biased peer observations in the collusion stress test, but cannot create Sybil identities, tamper with the Intelligent Decision Agent, or poison the offline training data.

**Design implications.** Under this threat model, a robust trust mechanism must satisfy four requirements: (R1) no single trust dimension can dominate the aggregate score, otherwise an adversary focusing on that dimension evades detection; (R2) trust weights must depend on observed context, because the adversary’s choice of which dimensions to attack is unknown a priori and may change over time; (R3) temporal consistency must be enforced, because dormant and on-off attackers exploit stale trust scores; and (R4) peer evidence must be down-weighted when an agent’s historical estimates become unreliable. MTIM’s cross-factor consistency check, policy-generated trust weighting, temporal decay, and Intelligent Decision Agent confidence weighting directly address R1–R4, respectively. The resulting protection is a continuous verification mechanism, not a complete replacement for cryptographic identity management or Sybil-resistant admission control.

### 3.4 Multi-factor Trust Evaluation Mechanism

MTIM assesses UAV credibility through five complementary trust dimensions, each capturing a distinct behavioral modality. The design principle is that a malicious node may evade detection in any single dimension by mimicking honest behavior, but sustaining consistent deception across five independent evidence channels is exponentially harder. We first define each dimension, then present the aggregation and cross-validation mechanisms.

#### 3.4.1 Trust Factor Decomposition

The five trust factors are selected to cover the minimum set of evidence channels needed by the threat model rather than to maximize the number of observable features. T1 captures behavioral consistency and exposes route or coordination deviations; T2 captures communication reliability and exposes selective forwarding or link manipulation; T3 captures resource-consumption anomalies that may indicate jamming, idle evasion, or compromised operation; T4 captures mission-facing task quality, which is the most direct evidence of whether a UAV contributes useful work; and T5 captures temporal reputation, which is necessary for dormant and on-off attackers that exploit one-shot observations. This choice creates computational overhead compared with single-factor scoring, but each factor corresponds to a distinct attack surface and the ablation in Table 6 evaluates the marginal value of the factor groups.

As illustrated in Figure 1, the trust evaluation is decomposed into multiple trust branches, each corresponding to a distinct trust factor. Let the overall trust score of UAV node  $i$  at time  $t$  be denoted as  $T_i(t) \in [0, 10]$ , which is a function of five fundamental trust dimensions (the formal aggregation mechanism is presented in Section 3.4.2), where each trust branch  $T_i^{(m)}(t) \in [0, 1]$  is defined as follows:

- **T1: Behavioral Trust**  $T_i^{(1)}(t)$ . This factor evaluates operational consistency, including adherence to flight paths, communication regularity, and response reliability:

$$T_i^{(1)}(t) = \frac{1}{N_b} \sum_{j=1}^{N_b} \mathbb{I}(a_{i,j}(t) = a_{i,j}^{expected}(t)) \quad (3)$$

where  $\mathbb{I}(\cdot)$  is the indicator function,  $a_{i,j}(t)$  represents the actual action in behavior category  $j$ , and  $a_{i,j}^{expected}(t)$  is the expected behavior. This yields  $T_i^{(1)}(t) \in [0, 1]$ .

- **T2: Communication Trust**  $T_i^{(2)}(t)$ . This factor assesses link quality and stability through packet delivery ratio, signal strength, and latency:

$$T_i^{(2)}(t) = w_p \cdot PDR_i(t) + w_s \cdot \frac{RSSI_i(t)}{RSSI_{max}} + w_l \cdot \left(1 - \frac{L_i(t)}{L_{max}}\right) \quad (4)$$

where  $PDR_i(t)$ ,  $RSSI_i(t)/RSSI_{max}$ , and  $1 - L_i(t)/L_{max}$  are each normalized to  $[0, 1]$ , and  $w_p = 0.4$ ,  $w_s = 0.3$ ,  $w_l = 0.3$  are weighting coefficients with  $w_p + w_s + w_l = 1$ . Thus  $T_i^{(2)}(t) \in [0, 1]$ .

- **T3: Energy Trust**  $T_i^{(3)}(t)$ . This factor monitors energy-consumption patterns and remaining battery life to detect anomalies indicative of malfunction or compromise:

$$T_i^{(3)}(t) = \frac{E_{remaining}(t)}{E_{initial}} \cdot \exp(-\alpha \cdot |C_i(t) - \bar{C}(t)|) \quad (5)$$

where  $\frac{E_{remaining}(t)}{E_{initial}} \in [0, 1]$  and the exponential term with  $\alpha = 0.5$  is also bounded in  $[0, 1]$ , yielding  $T_i^{(3)}(t) \in [0, 1]$ .

- **T4: Task Execution Trust**  $T_i^{(4)}(t)$ . This factor measures the success rate and quality of task completion, such as sensing, data relaying, or payload delivery:

$$T_i^{(4)}(t) = \frac{1}{N_t} \sum_{k=1}^{N_t} \left( \frac{Q_{i,k}(t)}{Q_k^{max}} \cdot \mathbb{I}(\text{task\_success}) \right) \quad (6)$$

where each term is normalized to  $[0, 1]$ , so  $T_i^{(4)}(t) \in [0, 1]$ .

- **T5: Historical Trust**  $T_i^{(5)}(t)$ . This factor incorporates long-term reputation scores and past behavior records to preserve temporal continuity in trust evaluation:

$$T_i^{(5)}(t) = \beta \cdot T_i^{(5)}(t-1) + (1-\beta) \cdot T_i^{inst}(t) \quad (7)$$

where both the historical and instantaneous components are in  $[0, 1]$ , and the smoothing factor  $\beta = 0.7$ . Thus  $T_i^{(5)}(t) \in [0, 1]$ .

### 3.4.2 Multi-factor Trust Aggregation

Each trust branch independently computes a sub-score  $T_i^{(m)}(t) \in [0, 1]$  based on real-time data streams and historical logs. These sub-scores are then aggregated into a global trust score  $T_i(t) \in [0, 10]$  through a weighted fusion mechanism:

$$T_i(t) = 10 \cdot \sum_{m=1}^5 \omega_m(t) \cdot T_i^{(m)}(t) \quad (8)$$

where  $\omega_m(t)$  are dynamic weighting coefficients satisfying  $\sum_{m=1}^5 \omega_m(t) = 1$ , and the factor of 10 maps the aggregated score to the  $[0, 10]$  range consistent with the classification thresholds  $\theta_{high} = 7.5$  and  $\theta_{low} = 3.0$ . These weights are produced by the MARL-style trust policy introduced in the next section, allowing the system to emphasize more reliable trust factors under varying operational conditions.

### 3.4.3 Cross-Validation Mechanism

The multi-factor design ensures that no single point of failure or manipulation can compromise the overall trust decision. To enhance robustness, a **cross-validation mechanism** is implemented between trust branches. For any pair of trust factors  $m$  and  $n$ , the consistency metric is defined as:

$$\Gamma_{m,n}(t) = \left| T_i^{(m)}(t) - T_i^{(n)}(t) \right| \quad (9)$$

If  $\Gamma_{m,n}(t) > \gamma_{th}$ , where the cross-validation threshold  $\gamma_{th} = 0.3$  (determined via sensitivity analysis over  $\{0.1, 0.2, 0.3, 0.4, 0.5\}$  with 0.3 yielding the best F-score on a validation subset), a potential inconsistency is flagged, triggering deeper inspection. For instance, even if a UAV exhibits normal communication behavior ( $T_i^{(2)}(t)$  high), but shows abnormal energy depletion ( $T_i^{(3)}(t)$  low), the system can still flag it as potentially compromised. This layered verification keeps trust continuously evaluated rather than assumed from any single evidence channel.

Furthermore, the framework supports **anomaly detection** through statistical analysis of trust factor deviations:

$$\mathcal{A}_i(t) = \mathbb{I} \left( \sum_{m=1}^5 \left( \frac{T_i^{(m)}(t) - \mu_m(t)}{\sigma_m(t)} \right)^2 > \chi_{\alpha,5}^2 \right) \quad (10)$$

where  $\mu_m(t)$  and  $\sigma_m(t)$  are the mean and standard deviation of trust factor  $m$  across the swarm,  $\chi_{\alpha,5}^2$  is the chi-squared threshold with 5 degrees of freedom at significance level  $\alpha = 0.05$ . This redundancy and diversity in trust evidence form the foundation for robust and adaptive trust management in UAV swarms, enabling the detection of sophisticated attacks that may manipulate multiple dimensions simultaneously.

### 3.5 MARL-Supported Incentive Policy and Multi-Agent Execution

Having defined the trust evidence channels, we now formulate how MTIM uses the resulting trust state to choose reward and penalty actions. MTIM employs  $M$  Thinker Agents, where each agent supervises a subset  $\mathcal{N}_j \subset \mathcal{U}$  of UAV nodes, maps local trust evidence to risk-aware trust-factor weights, and issues incentive actions that shape task allocation. The multi-agent design is used for local supervision under partial observability: each Thinker Agent sees only part of the swarm, while the Intelligent Decision Agent fuses their estimates into a global decision. The closed-loop system is modeled as a Dec-POMDP and follows a MARL design pattern, but the implementation uses an *offline-trained parameterized MARL trust-incentive surrogate*: the policy is trained from simulator-labeled trust streams and then frozen for all held-out evaluation seeds. We therefore avoid claiming online TD convergence or deployment-time reinforcement learning in the present study.

#### 3.5.1 Dec-POMDP Elements

We consider a UAV swarm consisting of  $N$  UAV nodes. The trust management system employs  $M$  Thinker Agents, where each agent  $j \in \{1, 2, \dots, M\}$  is responsible for a subset of UAV nodes. The overall system can be modeled as a Dec-POMDP (decentralized partially observable Markov decision process) [37]:

$$\mathcal{G} = \langle \mathcal{S}, \mathcal{O}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle \quad (11)$$

where  $\mathcal{S}$  is the global state space,  $\mathcal{O}$  is the joint observation space,  $\mathcal{A}$  is the joint action space,  $\mathcal{P}$  is the state transition probability,  $\mathcal{R}$  is the reward function, and  $\gamma \in [0, 1]$  is the discount factor.

**State Space:** The state  $s \in \mathcal{S}$  at time  $t$  is defined as the tuple of all UAV nodes' multi-factor trust scores:

$$s(t) = (\mathbf{T}_1(t), \mathbf{T}_2(t), \dots, \mathbf{T}_N(t)),$$

where  $\mathbf{T}_k(t) = [T_k^{(1)}(t), T_k^{(2)}(t), T_k^{(3)}(t), T_k^{(4)}(t), T_k^{(5)}(t)]$  denotes the five-dimensional trust vector for node  $k$ , with each component in  $[0, 1]$ . The corresponding global trust score  $\mathcal{T}_k(t) \in [0, 10]$  is computed by the Intelligent Decision Agent as described in Section 3.5.3.

**Observation Space:** Each Thinker Agent  $j$  observes a partial view of the global state, consisting only of the local trust sub-scores for the nodes in its supervised subset  $\mathcal{N}_j$ :  $o_j(t) = (\mathbf{T}_k(t))_{k \in \mathcal{N}_j} \cup (c_{jk}(t))_{k \in \mathcal{N}_j}$ , where  $c_{jk}(t)$  is the confidence score assigned by agent  $j$  to its own assessment of node  $k$ .

**Action Space:** The action  $a_j(t) \in \mathcal{A}_j$  for Thinker Agent  $j$  is a hybrid parameterized action consisting of (1) a continuous trust-weight vector  $\omega_j(t) = [\omega_j^{(1)}(t), \dots, \omega_j^{(5)}(t)] \in \Delta^4$  on the probability simplex, and (2) a discrete incentive allocation vector  $\pi_j(t) = [\pi_{jk}(t)]_{k \in \mathcal{N}_j}$  specifying the incentive level for each supervised node. Incentive levels are quantized into five bins  $\{-2, -1, 0, 1, 2\}$ , corresponding to strong penalty, weak penalty, neutral allocation, weak reward, and strong reward. Thus, trust evaluation provides the state used to judge behavior, while the learned policy decides whether the node should be rewarded with more favorable task participation or penalized through restricted access. If a sampled incentive vector violates the cycle budget  $B_{\text{total}}(t)$ , incentives are projected greedily by retaining actions with the largest expected utility gain per unit cost until the budget is satisfied.

**Architecture:** The implemented policy follows centralized offline training and decentralized execution, matching the CTDE motivation of MARL while using a supervised surrogate loss. During training, the simulator provides node labels only for the loss function; labels are never included in the deployed observations. During execution, each Thinker Agent acts only from local trust evidence and fixed learned parameters, ensuring that all reported decisions are label-free.

**State Transition:** The state transition  $\mathcal{P}$  is defined implicitly through the trust update equations (Eqs. 25–26) in Section 3.5.4, which govern how trust scores evolve based on observations, incentives, and temporal decay.

### 3.5.2 Executable Parameterized MARL Trust-Incentive Network

The implemented policy network uses per-node diagnostic features extracted from the post-warm-up trust streams. These features include recent weakness in T1, T2, T4, and T5; low-quantile T1/T4 evidence; low-score duration; T1/T4 amplitude and standard deviation; positive degradation slope; and directional cross-factor gaps. The same feature mapping is used by MTIM and by the learned baselines introduced in Section 4.1.

For node  $k$ , let  $\mathbf{x}_k \in \mathbb{R}^{13}$  denote the standardized diagnostic feature vector. A shared encoder maps  $\mathbf{x}_k$  to a hidden representation  $h_k$ :

$$h_k = f_\theta(\mathbf{x}_k). \quad (12)$$

Four output heads are then computed:

$$\begin{aligned} p_k &= \sigma(g_r(h_k)), \\ \omega_k^{(m)} &= \frac{\exp(g_\omega^{(m)}(h_k))}{\sum_{\ell=1}^5 \exp(g_\omega^{(\ell)}(h_k))}, \\ \mathbf{q}_k^\pi &= g_\pi(h_k), \\ v_k &= g_v(h_k), \end{aligned} \quad (13)$$

where  $p_k$  is the malicious-risk probability,  $\omega_k$  is the trust-factor weight vector,  $\mathbf{q}_k^\pi$  contains the five incentive-bin scores, and  $v_k$  is an auxiliary value estimate. The local diagnostic trust score is

$$\hat{T}_{jk} = 10(1 - p_k), \quad (14)$$

and a node is flagged as malicious when  $p_k > 0.60$ . The selected incentive bin is

$$\pi_{jk}^{raw} = \arg \max_{\ell \in \{-2, -1, 0, 1, 2\}} q_{k,\ell}^\pi, \quad (15)$$

followed by the budget projection described above.

### 3.5.3 Offline MARL Surrogate Training Objective

The executable training objective is a MARL surrogate objective with supervised and auxiliary terms rather than an online temporal-difference update. Let  $y_k \in \{0, 1\}$  be the simulator label used only during offline training. Let  $\omega_k^*$  be a heuristic trust-weight target derived from the same diagnostic evidence,  $\pi_k^*$  an incentive-bin target derived from the label and risk hint, and  $v_k^*$  a bounded service-value target. The policy minimizes

$$\begin{aligned} \mathcal{L}(\theta) &= \text{BCE}_w(y_k, p_k) + 0.25 \|\omega_k - \omega_k^*\|_2^2 \\ &\quad + 0.20 \text{CE}(\pi_k^*, \mathbf{q}_k^\pi) + 0.10 (v_k - v_k^*)^2. \end{aligned} \quad (16)$$

The primary risk term is weighted binary cross-entropy with a positive-class weight of 2.0. Training uses seeds 0–9 and a fixed decision threshold of 0.60. Evaluation seeds 1000–1049 are not used for training, threshold selection, or auxiliary-target tuning. This implementation makes the method-code correspondence explicit: MTIM’s reported policy is an offline-trained MARL-style trust-incentive policy, not a deployment-time online reinforcement-learning implementation.

The Intelligent Decision Agent in Figure 1 serves as the decision node that combines verified trust-branch outputs and Thinker-Agent estimates to produce the total trust score and incentive policy. This introduces a coordination point, so the claim is not that MTIM eliminates coordination cost. Instead, the hierarchical design addresses partial observability: local Thinker Agents observe different UAV subsets, while the Intelligent Decision Agent reconciles their estimates using confidence-aware uncertainty quantification (Eq. 19). The decision node can be replicated with Byzantine fault-tolerant consensus in safety-critical deployments, though we leave this extension to future work. The Intelligent Decision Agent uses the following confidence-weighted inference mechanism:

**Total Trust Score Aggregation:** The Intelligent Decision Agent computes the total trust score  $\mathcal{T}_k(t)$  for each UAV node  $k$  by combining the individual trust scores from all Thinker Agents that have observed or interacted with node  $k$ :

$$\mathcal{T}_k(t) = \frac{\sum_{j=1}^M \phi_{jk}(t) \cdot \hat{T}_{jk}(t)}{\sum_{j=1}^M \phi_{jk}(t)} \quad (17)$$

where  $\phi_{jk}(t) \in [0, 1]$  is the confidence weight assigned to Thinker Agent  $j$ 's assessment of node  $k$ , computed based on the agent's historical accuracy:

$$\phi_{jk}(t) = \frac{\exp(-\gamma \cdot \varepsilon_{jk}(t))}{\sum_{j'=1}^M \exp(-\gamma \cdot \varepsilon_{j'k}(t))} \quad (18)$$

with  $\varepsilon_{jk}(t) = \frac{1}{t} \sum_{\tau=1}^t \left| \hat{T}_{jk}(\tau) - T_{jk}^{obs}(\tau) \right|$  representing the mean absolute error of agent  $j$  for node  $k$  over time, and  $\gamma = 2.0$  a scaling parameter.

**Uncertainty Quantification:** The Intelligent Decision Agent also quantifies the uncertainty in global trust assessments using a Bayesian posterior variance:

$$\sigma_k^2(t) = \frac{\sum_{j=1}^M \phi_{jk}(t) \left( \hat{T}_{jk}(t) - \mathcal{T}_k(t) \right)^2}{\sum_{j=1}^M \phi_{jk}(t)} + \sigma_0^2 \quad (19)$$

where  $\sigma_0^2 = 0.1$  represents the inherent measurement noise variance. This uncertainty metric is crucial for risk-aware decision-making in zero-trust environments.

**Final Trust Classification:** Based on the total trust score and uncertainty, the Intelligent Decision Agent classifies nodes into trust categories:

$$C_k(t) = \begin{cases} \text{Trusted} & \text{if } \mathcal{T}_k(t) \geq \theta_{\text{high}} \text{ and } \sigma_k(t) \leq \sigma_{\text{th}} \\ \text{Suspicious} & \text{if } \theta_{\text{low}} \leq \mathcal{T}_k(t) < \theta_{\text{high}} \text{ or } \sigma_k(t) > \sigma_{\text{th}} \\ \text{Malicious} & \text{if } \mathcal{T}_k(t) < \theta_{\text{low}} \text{ and } \sigma_k(t) \leq \sigma_{\text{th}} \end{cases} \quad (20)$$

where  $\theta_{\text{high}}$  and  $\theta_{\text{low}}$  are trust thresholds, and  $\sigma_{\text{th}}$  is the uncertainty threshold.

The Intelligent Decision Agent generates the incentive policy that coordinates the actions of individual Thinker Agents to promote network-wide trustworthiness. The incentive policy  $\Pi(t)$  is a mapping from global trust assessments to resource allocation decisions:

$$\Pi(t) = \{ \pi_{jk}(t) : j \in [1, M], k \in \mathcal{N}_j \}, \quad (21)$$

where  $\pi_{jk}(t) \in \mathcal{A}_{\text{inc}} = \{-2, -1, 0, 1, 2\}$  is the incentive action for node  $k$  by agent  $j$ , encoding strong penalty ( $-2$ ) through strong reward ( $+2$ ). The policy is determined by solving a constrained optimization that maximizes expected network-wide utility:

$$\max_{\Pi(t)} \mathbb{E} \left[ \sum_{k=1}^N U_k(\mathcal{T}_k(t), \pi_{\cdot k}(t)) \right], \quad (22)$$

subject to the budget constraint  $\sum_{j=1}^M \sum_{k \in \mathcal{N}_j} c(\pi_{jk}(t)) \leq B_{\text{total}}(t)$ , where  $c(\pi) = |\pi|$  quantifies the resource cost of each incentive action. The per-node utility function is:

$$U_k(\mathcal{T}_k(t), \pi_{\cdot k}(t)) = \mathcal{T}_k(t) \cdot (1 + \delta \cdot \mathbb{I}(\pi_{\cdot k}(t) > 0)) - \kappa \cdot \mathbb{I}(\mathcal{T}_k(t) < \theta_{\text{low}}), \quad (23)$$

with  $\delta = 0.2$  and  $\kappa = 2.0$ . The first term rewards high-trust nodes with proportional benefits, while the second term imposes a fixed penalty on nodes classified as Malicious, creating a natural pressure toward trustworthy behavior. In practice, the optimization in Eq. 22 is solved greedily per update cycle due to the 50 s incentive interval constraint (Table 2).

### 3.5.4 State Update Mechanism for Trust Evolution

The state update mechanism governs how trust scores evolve over time as a function of node behavior, incentive actions, and network conditions. Trust evaluation uses Bayesian inference, and the overall trust dynamics constitute a Markov process because the trust score at time  $t + 1$  depends only on the trust score at time  $t$  and the current observations and incentives, not on the full history. The state transition can be decomposed as:

$$\mathbf{T}(t + 1) = f(\mathbf{T}(t), \mathbf{A}^{\text{inc}}(t), \mathbf{A}^{\text{task}}(t), \mathbf{E}(t)) \quad (24)$$

where  $\mathbf{T}(t) \in \mathbb{R}^N$  is the vector of global trust scores,  $\mathbf{A}^{\text{inc}}(t)$  is the incentive action matrix,  $\mathbf{A}^{\text{task}}(t)$  captures task execution outcomes, and  $\mathbf{E}(t)$  represents environmental factors. The trust score for each node  $k$  evolves according to the following three-stage process:

**Temporal Trust Decay:** To prevent trust stagnation and enforce continuous verification, we incorporate a temporal decay factor:

$$\tilde{\mathcal{T}}_k(t+1) = \mathcal{T}_k(t) \cdot \exp(-\rho \cdot \Delta t) \quad (25)$$

where  $\rho = 0.01$  is the decay rate (see Table 2) and  $\Delta t = 10$  s is the trust update interval. This ensures that trust gradually diminishes without positive reinforcement.

**Trust Update Based on Observations:** The final trust score is updated via a Bayesian filter that combines the decayed prior with new observations:

$$\begin{aligned} \mathcal{T}_k(t+1) &= a_k(t) \tilde{\mathcal{T}}_k(t+1) + (1 - a_k(t)) \\ &\quad \cdot T_k^{obs}(t+1), \\ a_k(t) &= \frac{\sigma_{obs}^2}{\sigma_k^2(t) + \sigma_{obs}^2}. \end{aligned} \quad (26)$$

where  $\sigma_{obs}^2 = 0.5$  is the observation noise variance, and  $T_k^{obs}(t+1)$  is the aggregated observed trust score from all Thinker Agents. This form assigns more weight to new observations when the prior uncertainty  $\sigma_k^2(t)$  is high, and more weight to the decayed prior when the current estimate is confident. The observation term is:

$$T_k^{obs}(t+1) = \frac{1}{|\mathcal{M}_k(t+1)|} \sum_{j \in \mathcal{M}_k(t+1)} T_{jk}^{obs}(t+1) \quad (27)$$

with  $\mathcal{M}_k(t+1)$  denoting the set of agents that observed node  $k$  at time  $t+1$ .

**Boundedness and Empirical Stability:** The trust update mechanism in Eqs. 25–26 admits bounded updates under the  $[0, 10]$  trust scale. Since  $\mathcal{T}_k(t) \in [0, 10]$  and  $\tilde{\mathcal{T}}_k(t+1) = \mathcal{T}_k(t) \cdot \exp(-\rho \Delta t) \in [0, 10]$ , and the Bayesian fusion in Eq. 26 is a convex combination of two bounded terms (with coefficients in  $[0, 1]$  that sum to one), the trust score remains in  $[0, 10]$  for all  $t$ , preventing numerical divergence.

We do not claim a finite-sample convergence guarantee for the full Dec-POMDP, because the empirical policy is trained offline rather than by solving the Dec-POMDP online. Instead, we provide boundedness of the trust-state update and empirical stability evidence. Section 4.2 shows that trust scores reach quasi-steady-state separation within the 5000 s simulation window for all tested configurations.

**Trust-Driven Network Evolution:** The state update mechanism creates a positive feedback loop where trustworthy behavior is rewarded with higher trust scores and corresponding incentives, while malicious behavior leads to trust degradation and restricted access. Over time, this mechanism induces a stabilization effect where the distribution of trust scores across the swarm converges, as verified empirically in Section 4.2.

### 3.5.5 Evaluation Metrics

This section defines the main evaluation metrics used throughout the experiments. All metrics are computed at the end of each simulation run (after the 1000 s warm-up period).

**Trust Classification Metrics:** For malicious-node detection, the confusion matrix is defined over the three-class output of Eq. 20, with Suspicious mapped to Malicious for binary evaluation. Let TP, TN, FP, FN denote true/false positives/negatives, respectively. Detection accuracy, precision, recall, and  $F_1$ -score are:


$$\begin{aligned} \text{Accuracy} &= \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}, & \text{Precision} &= \frac{\text{TP}}{\text{TP} + \text{FP}}, \\ \text{Recall} &= \frac{\text{TP}}{\text{TP} + \text{FN}}, & F_1 &= \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}. \end{aligned} \quad (28)$$

Ground-truth labels are assigned at simulation initialization according to the attack model (Section 3.3) and are never observable by the Thinker Agents during execution.

**Mission-Level Metric:** The task success rate (TSR) measures mission continuity independently of detection accuracy:

$$\text{TSR} = N_{\text{succ}} / N_{\text{all}}, \quad (29)$$

where a task succeeds iff it completes before the mission timeout with at least one Trusted UAV contributing, and fails if it times out or is assigned exclusively to Suspicious/Malicious nodes.



**Figure 2:** MTIM state-update and incentive-feedback mechanism aligned with Figure 1. The forward path maps Summarization-Agent observations and verified trust branches to Thinker-Agent policy actions, Intelligent Decision Agent confidence-weighted scoring, temporally decayed status updates, and classification-driven reward-penalty decisions. The feedback path returns incentive outcomes to the next update interval so that task allocation and future trust evolution are coupled.

**Efficiency Metrics:** Communication overhead  $C_{\text{comm}}$  (bytes/UAV/update cycle) and per-cycle evaluation latency  $T_{\text{eval}}$  (ms) are reported to quantify deployment feasibility. The evaluation latency is measured end-to-end from observation ingestion to Intelligent Decision Agent classification output.

### 3.5.6 Algorithm Summary

Table 1 consolidates the MTIM trust evaluation and incentive loop for a single update cycle.

**Table 1:** MTIM trust evaluation and incentive cycle for one update interval  $\Delta t$ .

|                |                                                                                                                                                      |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Input:</b>  | Current trust vectors $\{\mathbf{T}_k(t)\}_{k=1}^N$ and agent assignments $\{\mathcal{N}_j\}_{j=1}^M$ .                                              |
| <b>Output:</b> | Global trust scores $\{\mathcal{T}_k(t)\}_{k=1}^N$ , classifications $\{C_k(t)\}_{k=1}^N$ , and incentives $\{\pi_{jk}(t)\}$ .                       |
| 1              | For each Thinker Agent $j = 1, \dots, M$ , collect local observations $o_j(t) = \{\mathbf{T}_k(t)\}_{k \in \mathcal{N}_j}$ .                         |
| 2              | Extract diagnostic features and compute $p_k$ , $\omega_j(t)$ , and $\pi_j(t)$ using the frozen policy in Eq. 13.                                    |
| 3              | Estimate local trust as $\hat{T}_{jk}(t) = 10(1 - p_k)$ and retain $\omega_j(t)$ for factor-level interpretability.                                  |
| 4              | Apply the cross-validation check in Eq. 9 and flag inconsistent nodes.                                                                               |
| 5              | Update Intelligent Decision Agent confidence weights $\phi_{jk}(t)$ via Eq. 18.                                                                      |
| 6              | Compute global trust $\mathcal{T}_k(t)$ and uncertainty $\sigma_k^2(t)$ using Eqs. 17–19.                                                            |
| 7              | Classify nodes with Eq. 20.                                                                                                                          |
| 8              | Solve the constrained incentive allocation in Eq. 22 greedily under $B_{\text{total}}(t)$ .                                                          |
| 9              | Apply temporal decay and Bayesian update using Eqs. 25–26.                                                                                           |
| 10             | Log observations and decisions for audit; deployment-time policy parameters remain fixed unless an offline retraining cycle is explicitly triggered. |

## 4 Performance Evaluation

Our evaluation targets four questions:

(Q1) **Detection performance:** How accurately does MTIM classify nodes under varying malicious-node ratios and attack patterns, and how does it compare against state-of-the-art baselines?

- (Q2) **Incentive effect:** Does coupling trust scores to reward and penalty decisions improve task success beyond risk-only classification?
- (Q3) **Trust-factor contribution:** Which evidence factors and temporal mechanisms are necessary for robust behavior assessment?
- (Q4) **Scalability:** How do latency, overhead, and accuracy scale with swarm size?

## 4.1 Simulation Setup

**Hardware and software.** The evaluation was implemented with a reproducible trust-stream simulator and an offline-trained policy network. All compared methods consume the same normalized T1–T5 evidence streams and are evaluated under identical random seeds and attack schedules. A packet-trace conversion check is included as an interface sanity test; the numerical results reported below, however, should be read as trust-layer simulation evidence rather than multi-seed packet-level channel validation or flight-test validation.

**Network and mobility.** All scenario parameters are consolidated in Table 2. The reference simulator abstracts mobility and channel variation into seeded trust-factor streams that include node heterogeneity, intermittent observation gaps, and attack-dependent perturbations. The mission duration is 5000 s; the first 1000 s serve as a warm-up during which trust scores are accumulated but no incentive actions are enforced.

**Adversary configuration.** Malicious nodes implement the three attack strategies defined in Section 3.3: constant (adversarial from  $t = 0$ ), dormant (switch to adversarial at  $t = T_{\text{warm}} + \Delta$ , where  $\Delta \sim \mathcal{U}(200, 500)$  s), and on-off (toggle with period  $T_{\text{toggle}} = 500$  s). In mixed-attack experiments, each type accounts for one-third of the malicious population. The malicious-node ratio  $|\mathcal{M}|/N$  is varied from 5% to 50% in anti-manipulation experiments and fixed at 30% for the main comparison.

**Training protocol.** MTIM uses a compact parameterized MARL trust–incentive surrogate matching the action structure in Section 3.5: a shared encoder outputs a malicious-risk logit, a five-dimensional trust-weight softmax, five incentive-bin scores, and an auxiliary value estimate. Training seeds 0–9 fit the network using trust-evidence features derived from recent weakness, low-quantile T1/T4 evidence, low-score duration, amplitude, degradation slope, and directional cross-factor gaps. The risk head supports trust assessment with weighted binary cross-entropy, while the incentive head learns reward/penalty bins from label-derived risk hints and service-value targets. The decision threshold is fixed at 0.60 before held-out evaluation. Evaluation seeds 1000–1049 are never used for training or threshold selection.

**Evaluation protocol.** Each configuration is evaluated over 50 independent trials with different random seeds controlling malicious-node identities, attack schedules, trust-factor noise, and observation gaps. Unless otherwise specified, the evaluation seeds are indexed from 1000 to 1049. All reported results are computed from the trained execution policy and fixed baseline thresholds, and are reported as mean  $\pm$  standard deviation.

**Baselines.** We compare MTIM against six methods spanning the design spectrum from naive fusion to learned trust classification:

- **MFA (Multi-factor Average):** Equal-weighted average of T1–T5, representing the performance floor achievable without any learning or adaptation. Included to isolate the gain attributable to adaptive weighting.
- **FUBA [8]:** Fuzzy logic-based trust analytics for FANETs, reimplemented as membership-function trust fusion over the shared T1–T5 stream.
- **GALTrust [10]:** Generative-adversarial trust management, reimplemented as a distance-to-honest-profile anomaly score with volatility and consistency terms over the same stream.
- **GDM-DTM [23]:** Group decision-making dynamic trust management, reimplemented as weighted group-decision trust fusion with a consistency penalty.

**Table 2:** Reference-simulation parameters. Physical-layer and mobility values are scenario descriptors; the evaluation operates at the trust-factor-stream level.

| Category | Parameter                                  | Value                                                   |
|----------|--------------------------------------------|---------------------------------------------------------|
| Scenario | Simulation area                            | $10 \times 10 \text{ km}^2$                             |
|          | Mission duration                           | 5000 s (warm-up: 1000 s)                                |
|          | UAV count $N$                              | 5–50 (default: 15)                                      |
| Mobility | UAV speed                                  | $\mathcal{U}(3, 15) \text{ m/s}$                        |
|          | Communication range $R_{\text{comm}}$      | 300 m                                                   |
|          | Inter-UAV bandwidth                        | 5 Mbit/s                                                |
| Energy   | Initial energy                             | $\mathcal{U}(50, 200) \text{ Wh}$                       |
|          | Power consumption                          | $\mathcal{U}(100, 120) \text{ W}$                       |
| Trust    | Trust update interval $\Delta t$           | 10 s                                                    |
|          | Incentive update interval                  | 50 s                                                    |
|          | Classification thresholds                  | $\theta_{\text{high}} = 7.5, \theta_{\text{low}} = 3.0$ |
|          | Uncertainty threshold $\sigma_{\text{th}}$ | 0.5                                                     |
| Training | Policy network                             | PyTorch parameterized MARL trust–incentive surrogate    |
|          | Training seeds                             | 0–9                                                     |
|          | Evaluation seeds                           | 1000–1049                                               |
|          | Decision threshold                         | 0.60 (fixed before held-out evaluation)                 |
|          | Python dependencies                        | numpy, pandas, matplotlib, scipy, torch                 |

- **LR-Trust:** A logistic risk classifier trained on exactly the same diagnostic features and training seeds as MTIM, but without MTIM’s trust-weight, incentive, or value heads.
- **MLP-Trust:** A two-layer neural risk classifier trained on exactly the same diagnostic features and training seeds as MTIM, but without MTIM’s trust-weight, incentive, or value heads.

All baselines share the identical generated trust-factor streams and ground-truth labels as MTIM. Because official source code and original feature pipelines for FUBA, GALTrust, and GDM-DTM are not bundled, those baseline results should be interpreted as paper-level reimplementations under a common T1–T5 interface, not as definitive re-evaluations of the original authors’ implementations. LR-Trust and MLP-Trust are added specifically to test whether MTIM’s advantage persists against stronger learned classifiers with identical input information.

**Baseline fairness and statistical protocol.** To reduce implementation bias, all compared methods consume the same five normalized trust-factor streams and are evaluated on the same 50 held-out attack seeds. The MTIM, LR-Trust, and MLP-Trust networks are trained only on seeds 0–9 and are not re-tuned on evaluation seeds 1000–1049. The baseline thresholds are fixed before evaluation.

Statistical comparison is performed at the trial level. For each evaluation seed, we compute paired differences between MTIM and each baseline for Accuracy, F-score, and TSR. Paired  $t$ -test statistics use a normal approximation and Holm correction across the 18 comparisons (6 baselines  $\times$  3 metrics), together with paired differences and standardized effect sizes. MTIM significantly improves Accuracy, F-score, and TSR over MFA, FUBA, GALTrust, GDM-DTM, and LR-Trust after Holm correction at  $p < 0.01$ . Against MLP-Trust, Accuracy and F-score are not significantly different, while MTIM significantly improves TSR. Communication overhead and latency are treated as resource trade-offs rather than superiority objectives.

**Collusion stress test.** To test whether Intelligent Decision Agent confidence weighting can suppress biased peer evidence, we additionally evaluate a limited collusion setting in which 20% of malicious nodes submit inflated recommendations for other malicious nodes while continuing their assigned behavioral attack pattern. Because MTIM does not implement Sybil-resistant admission control, this stress test assumes fixed physical node identities and focuses only on false peer-report inflation. Results from this setting are reported qualitatively in the sensitivity analysis and are used to bound, rather than broaden, the security claim.

(\*) MFA = Multi-factor Average. TSR = Task Success Rate. Comm. Overhead = per-UAV communication overhead per update cycle. Eval. Time = trust evaluation latency per update interval. Under Holm correction, MTIM’s Accuracy and F-score are significant over MFA, FUBA, GALTrust, GDM-DTM, and LR-Trust at  $p < 0.01$ , but not over MLP-Trust.

**Table 3:** Performance comparison of MTIM with baseline methods under the unified experimental setting (30% malicious-node ratio, 15 UAVs, 5000 s simulation). All values are mean  $\pm$  standard deviation over 50 independent trials.

| Method        | Accuracy                | F-score                 | TSR                     | Comm. Overhead (KB/UAV) | Eval. Time (ms) |
|---------------|-------------------------|-------------------------|-------------------------|-------------------------|-----------------|
| MFA           | 80.3% $\pm$ 10.5%       | 69.4% $\pm$ 16.3%       | 79.7% $\pm$ 4.8%        | 11.3 $\pm$ 1.8          | 10.2 $\pm$ 4.5  |
| FUBA [8]      | 84.1% $\pm$ 9.2%        | 76.1% $\pm$ 14.6%       | 82.5% $\pm$ 4.2%        | 16.7 $\pm$ 1.7          | 17.1 $\pm$ 4.1  |
| GALTrust [10] | 76.1% $\pm$ 6.0%        | 41.7% $\pm$ 21.8%       | 75.7% $\pm$ 4.2%        | 31.4 $\pm$ 1.8          | 40.6 $\pm$ 4.0  |
| GDM-DTM [23]  | 85.1% $\pm$ 8.3%        | 77.3% $\pm$ 12.6%       | 81.3% $\pm$ 3.9%        | 19.4 $\pm$ 1.7          | 21.5 $\pm$ 3.7  |
| LR-Trust      | 88.1% $\pm$ 9.6%        | 81.5% $\pm$ 15.4%       | 87.6% $\pm$ 3.6%        | 21.3 $\pm$ 1.8          | 24.5 $\pm$ 3.7  |
| MLP-Trust     | <b>94.7%</b> $\pm$ 5.6% | <b>92.1%</b> $\pm$ 7.9% | <u>88.3%</u> $\pm$ 3.9% | 25.2 $\pm$ 2.0          | 36.1 $\pm$ 3.6  |
| MTIM (ours)   | <u>94.4%</u> $\pm$ 5.6% | <u>91.8%</u> $\pm$ 8.0% | <b>90.5%</b> $\pm$ 3.7% | 34.0 $\pm$ 1.8          | 58.3 $\pm$ 4.3  |

MTIM’s TSR is significant over all baselines, including MLP-Trust. Overhead and latency are reported as trade-offs rather than optimization objectives.

**Analysis of Table 3.** MTIM achieves 94.4% accuracy, 91.8% F-score, and 90.5% TSR. Three observations explain the result:

(i) **Learned temporal evidence vs. manual/profile-based fusion.** MFA reaches 80.3% accuracy, FUBA reaches 84.1%, and GDM-DTM reaches 85.1%. MTIM improves over GDM-DTM by 9.3 pp in accuracy and 14.5 pp in F-score because it uses low-quantile and degradation-slope evidence that exposes dormant and on-off attackers. This advantage persists at the 50% malicious-node ratio (95.5% vs. 82.4%; Table 4).

(ii) **Strong learned baselines narrow the detection claim.** LR-Trust improves over non-neural fusion baselines but still trails MTIM by 6.3 pp in accuracy. MLP-Trust reaches 94.7% accuracy and 92.1% F-score, which are statistically indistinguishable from MTIM’s 94.4% and 91.8%. Thus, the evidence does not support claiming that MTIM is a strictly better detector than a strong neural risk classifier with the same features; the stronger claim is that MTIM links comparable detection quality to better mission-level action.

(iii) **Incentive coupling closes the mission loop.** MTIM achieves the best TSR (90.5%), exceeding GDM-DTM by 9.2 pp and MLP-Trust by 2.2 pp. The MLP-Trust comparison is important: detection alone is nearly tied, but MTIM’s incentive head and task-allocation coupling produce a measurable mission-level gain.

#### Overhead trade-off.

MTIM’s per-cycle latency (58.3 ms) and communication overhead (34.0 KB/UAV) are higher than all baselines, including MLP-Trust (36.1 ms and 25.2 KB/UAV), but remain well within the 10 s update budget in the reference simulator. Most latency comes from multi-agent aggregation, diagnostic-feature extraction, and incentive-policy output. If a mission requires a shorter update interval, the number of Thinker Agents  $M$  can be reduced, with sparser observation coverage as the trade-off.

**Robustness across malicious-node ratios.** Table 4 reports accuracy and TSR at 10%, 30%, and 50% malicious-node ratios. MTIM retains the best TSR across all three ratios. For detection accuracy, MTIM and MLP-Trust are closely matched: MLP-Trust is slightly higher at 10% and 30%, whereas MTIM is higher at 50%. GALTrust has strong accuracy at 10% because its anomaly-profile threshold is conservative when the attack population is sparse, but it degrades sharply as the malicious ratio increases. MTIM’s accuracy remains between 94.4% and 96.8%, while TSR remains above 88%, indicating that the trained policy continues to support mission-level decisions under dense adversarial pressure. The slightly higher accuracy at 50% than at 30% reflects stronger aggregate separation in the generated attack streams and should not be interpreted as a universal monotonic robustness guarantee.

**Policy fitting stability.** On training seeds 0–9, the final MTIM policy reaches 100.0% training accuracy and F-score at threshold 0.60, with mean malicious risk 0.9997 and mean benign risk 0.0004. This complete training-set separation indicates that the generated training streams are highly learnable, so the reported performance is based on held-out seeds 1000–1049 rather than training metrics. The learned trust-weight head assigns the largest average weights to T4 (0.323) and T1 (0.314), followed by

**Table 4:** Detection accuracy and TSR at three malicious-node ratios ( $N = 15$ , 50 trials). Best values in **bold**, second-best underlined.

| Method      | 10% Malicious         |                       | 30% Malicious         |                       | 50% Malicious         |                       |
|-------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|
|             | Accuracy              | TSR                   | Accuracy              | TSR                   | Accuracy              | TSR                   |
| MFA         | 86.1 $\pm$ 9.5        | 87.4 $\pm$ 3.1        | 80.3 $\pm$ 10.5       | 79.7 $\pm$ 4.8        | 76.9 $\pm$ 8.2        | 71.1 $\pm$ 4.9        |
| FUBA        | 83.7 $\pm$ 9.8        | 87.8 $\pm$ 3.3        | 84.1 $\pm$ 9.2        | 82.5 $\pm$ 4.2        | 81.1 $\pm$ 9.7        | 73.8 $\pm$ 5.2        |
| GALTrust    | <u>91.2</u> $\pm$ 4.4 | 85.2 $\pm$ 3.2        | 76.1 $\pm$ 6.0        | 75.7 $\pm$ 4.2        | 63.7 $\pm$ 7.6        | 65.3 $\pm$ 3.8        |
| GDM-DTM     | 87.2 $\pm$ 7.9        | 88.5 $\pm$ 2.3        | 85.1 $\pm$ 8.3        | 81.3 $\pm$ 3.9        | 82.4 $\pm$ 7.8        | 73.9 $\pm$ 4.5        |
| LR-Trust    | 90.4 $\pm$ 7.4        | 90.2 $\pm$ 2.0        | 88.1 $\pm$ 9.6        | 87.6 $\pm$ 3.6        | 88.1 $\pm$ 8.1        | 85.3 $\pm$ 4.6        |
| MLP-Trust   | <b>97.3</b> $\pm$ 4.3 | <u>90.6</u> $\pm$ 2.0 | <b>94.7</b> $\pm$ 5.6 | <u>88.3</u> $\pm$ 3.9 | <u>94.4</u> $\pm$ 5.1 | <u>85.6</u> $\pm$ 5.3 |
| MTIM (ours) | <u>96.8</u> $\pm$ 4.5 | <b>92.6</b> $\pm$ 1.9 | <u>94.4</u> $\pm$ 5.6 | <b>90.5</b> $\pm$ 3.7 | <b>95.5</b> $\pm$ 4.4 | <b>88.8</b> $\pm$ 4.6 |

T2 (0.145), T5 (0.137), and T3 (0.081), matching the threat model in which sustained task/behavior degradation is harder for evasive attackers to hide than energy behavior.


## 4.2 Trust Evolution Analysis

Figure 3 tracks the MTIM diagnostic trust index of benign and malicious UAVs from the post-warm-up phase to 5000 s under  $M \in \{2, 5, 10\}$  Thinker Agents. The index is computed from the trained risk policy and maps lower risk to higher trust on a 0–10 scale. Three dynamics merit attention:

**Divergence rate vs. agent count.** Benign nodes remain mostly above 8.0 after warm-up, whereas malicious nodes decline as low-quantile and degradation-slope evidence accumulates. The  $M = 5$  configuration shows a smoother decline than  $M = 2$  and  $M = 10$ , reflecting the trade-off between sparse observation coverage and aggregation latency. The default  $M = 5$  balances these effects in the reference simulator.


**Attack-pattern sensitivity.** Dormant attackers initially resemble benign nodes but decline after the post-warm-up switch because the policy uses low T1/T4 quantiles and positive degradation slope, not only the final average score. On-off attackers oscillate because honest phases partially restore recent averages, but amplitude and low-score-duration features keep the diagnostic index below the benign trajectory.

**Steady-state behavior.** After approximately 3800 s, the malicious diagnostic index remains near or below the Suspicious/Malicious boundary in most runs, while benign nodes remain high despite observation gaps. The residual fluctuations reflect seeded trust-factor noise and on-off behavior rather than formal convergence of the full Dec-POMDP.



**Figure 3:** Temporal evolution of the MTIM diagnostic trust index for benign (solid) and malicious (dashed) UAVs under  $M = 2, 5, 10$  Thinker Agents after warm-up. The index is recomputed from the trained risk policy over sliding trust-factor windows; lower values indicate higher malicious risk.

### 4.3 Comparative Performance Analysis



**Figure 4:** Comparative evaluation results. (a) Detection accuracy vs. malicious-node ratio: MTIM remains high across 5–50% malicious ratios and is closely matched by MLP-Trust, while GALTrust degrades sharply under dense adversarial pressure. (b) MTIM diagnostic trust-index dynamics under benign, constant-attack, dormant-attack, and on-off behavior. (c) TSR comparison at a 30% malicious-node ratio using the same baselines as Table 3; MTIM yields the highest task success rate.

**Anti-manipulation under varying malicious ratios (Figure 4(a)).** MTIM maintains high accuracy for malicious ratios of 20% and above, with 94.4% accuracy at 30% and 95.5% at 50%. MLP-Trust is a strong detection competitor under the same features, confirming that the diagnostic feature set itself carries substantial discriminative information. GALTrust is highly accurate at 5–10% malicious ratios but falls to 63.7% at 50%, indicating that the profile-distance adapter becomes unstable when the attack population substantially shifts the empirical trust distribution. MTIM does not show monotonic degradation across this generated ratio grid; this is favorable for the bounded setting but should be interpreted cautiously because stronger collusion and adaptive poisoning are not modeled.

**Trust dynamics under attack patterns (Figure 4(b)).** Constant and dormant attackers decline below the benign diagnostic index as low T1/T4 evidence accumulates. Dormant attackers fall after their switch despite initially high scores, while on-off attackers oscillate because their honest phases restore recent averages. This behavior supports a conservative interpretation: MTIM identifies many on-off attackers as high risk, but their oscillatory trajectory remains the hardest case and drives the wide F-score variance in Table 3.

**Task success comparison (Figure 4(c)).** MTIM achieves the highest TSR (90.5%), followed by MLP-Trust (88.3%) and LR-Trust (87.6%). The TSR ranking does not exactly mirror the accuracy ranking because task success depends on which nodes are selected after classification, not only on the binary detection count. The incentive mechanism therefore matters most at the mission-decision layer, as confirmed by the ablation in Table 6.

**Collusion stress-test interpretation.** Under inflated peer recommendations, MTIM’s Intelligent Decision Agent confidence weights down-weight agents whose historical estimates diverge from subsequently observed behavior. This reduces the impact of biased reports, but it does not eliminate collusion risk when the colluding group controls most local observations for a node. Therefore, the reported robustness should be interpreted as resistance to limited false-report inflation under fixed identities, not as a complete solution to collusion or Sybil attacks.

**Stress cases not claimed.** We explicitly do not claim robustness against three stronger adversaries: (i) Sybil identity inflation, where one physical attacker creates many unlinkable UAV identities; (ii) majority local-observation collusion, where colluding nodes dominate all observations available for a target node; and (iii) Intelligent Decision Agent compromise or training-pipeline poisoning. These cases require identity binding, Byzantine-resilient aggregation, secure bootstrapping, or robust training defenses that are orthogonal to MTIM’s trust-scoring layer. The present evaluation should therefore be read as a fixed-identity behavioral-manipulation study with a limited false-report stress test.

**ns-3 trace-driven sanity check.** To connect the trust-stream interface to packet/mobility observations, we converted the bundled 500 s ns-3-derived trace into T1–T5 streams and evaluated all methods on the converted trace. This is a single seed-1000 trace with 15 UAVs, 49 update steps, 10 benign nodes,

three constant attackers, and two on-off attackers; it is not a multi-seed packet-level validation. Table 5 shows that MTIM correctly preserves all benign nodes but detects only two of five malicious nodes on this short trace. In this trace-only setting, MFA, FUBA, and GDM-DTM detect three malicious nodes. Therefore, the ns-3 trace path should be interpreted as an interface-validation sanity check and as evidence that packet-level validation can expose harder short-horizon behavior than the main trust-stream simulator.

**Table 5:** Trace-driven validation on the bundled 500 s ns-3-derived trust stream (single seed, 15 UAVs, 49 update steps).

| Method      | Accuracy | Precision | Recall | F-score | TP/TN/FP/FN | Mean score (B/M) |
|-------------|----------|-----------|--------|---------|-------------|------------------|
| MFA         | 86.7%    | 100.0%    | 60.0%  | 75.0%   | 3/10/0/2    | 0.900/0.619      |
| FUBA        | 86.7%    | 100.0%    | 60.0%  | 75.0%   | 3/10/0/2    | 0.939/0.574      |
| GALTrust    | 66.7%    | 50.0%     | 60.0%  | 54.5%   | 3/7/3/2     | 0.718/0.428      |
| GDM-DTM     | 86.7%    | 100.0%    | 60.0%  | 75.0%   | 3/10/0/2    | 0.938/0.586      |
| LR-Trust    | 80.0%    | 100.0%    | 40.0%  | 57.1%   | 2/10/0/3    | 1.000/0.600      |
| MLP-Trust   | 80.0%    | 100.0%    | 40.0%  | 57.1%   | 2/10/0/3    | 1.000/0.584      |
| MTIM (ours) | 80.0%    | 100.0%    | 40.0%  | 57.1%   | 2/10/0/3    | 1.000/0.482      |

#### 4.4 Ablation Study

To quantify the marginal contribution of each architectural component, we conduct a structured ablation study. Each variant removes exactly one component while keeping all others intact. All variants are evaluated under the identical setting (30% malicious,  $N = 15$ , 50 trials).

**Table 6:** Ablation study of MTIM components (30% malicious-node ratio, 15 UAVs). All values are mean  $\pm$  std over 50 trials.

| Variant                   | Accuracy          | F-score           | TSR              |
|---------------------------|-------------------|-------------------|------------------|
| MTIM (full)               | 94.4% $\pm$ 5.6%  | 91.8% $\pm$ 7.9%  | 90.5% $\pm$ 3.7% |
| w/o Cross-factor Checking | 92.9% $\pm$ 5.9%  | 89.5% $\pm$ 8.6%  | 89.5% $\pm$ 4.2% |
| w/o Bayesian Aggregation  | 84.1% $\pm$ 6.1%  | 67.7% $\pm$ 16.2% | 81.9% $\pm$ 4.1% |
| w/o Temporal Decay        | 79.5% $\pm$ 6.1%  | 52.8% $\pm$ 20.5% | 80.0% $\pm$ 3.8% |
| w/o Adaptive Policy       | 82.3% $\pm$ 7.0%  | 64.2% $\pm$ 18.2% | 80.9% $\pm$ 4.5% |
| w/o Incentive Mechanism   | 94.3% $\pm$ 5.7%  | 91.7% $\pm$ 8.0%  | 81.4% $\pm$ 3.7% |
| T1 Only                   | 80.8% $\pm$ 11.5% | 74.9% $\pm$ 14.2% | 82.6% $\pm$ 4.5% |
| T2 Only                   | 63.7% $\pm$ 11.7% | 53.3% $\pm$ 13.8% | 76.2% $\pm$ 5.1% |
| T3 Only                   | 57.6% $\pm$ 12.2% | 56.0% $\pm$ 9.8%  | 75.9% $\pm$ 6.1% |
| T4 Only                   | 82.3% $\pm$ 10.1% | 76.5% $\pm$ 12.9% | 83.2% $\pm$ 4.3% |
| T5 Only <sup>b</sup>      | 84.0% $\pm$ 9.7%  | 79.2% $\pm$ 13.1% | 82.6% $\pm$ 3.9% |

<sup>b</sup> T5-only operates on the raw T1–T4 instantaneous values as its input features and applies only exponential smoothing ( $\beta = 0.7$ ) without any aggregation or learning. It represents the best achievable performance using temporal smoothing alone.

**Analysis of Table 6.** The ablation results reveal a more nuanced component picture than the conceptual model alone suggests:

**Tier 1—Mission-critical: incentives and temporal/adaptive evidence.** Removing the incentive mechanism leaves detection accuracy essentially unchanged (94.3%) but reduces TSR from 90.5% to 81.4%, showing that the trust score must be coupled to task allocation to improve mission outcomes. Removing temporal decay reduces accuracy by 14.9 pp and F-score by 39.0 pp, confirming that dormant and on-off attacks cannot be handled reliably from final averages alone. Removing the adaptive policy

reduces accuracy by 12.1 pp, because fixed aggregation cannot exploit low-quantile and degradation-slope evidence.

**Tier 2—Decision-agent fusion: Bayesian confidence weighting.** Removing Bayesian aggregation reduces accuracy by 10.3 pp and TSR by 8.6 pp. This supports the need for confidence-aware fusion among local estimates, but the effect is smaller than the temporal-policy effect in the current implementation.

**Tier 3—Secondary diagnostic safeguard: cross-factor checking.** Removing the directional cross-factor gap features reduces accuracy from 94.4% to 92.9% and TSR from 90.5% to 89.5%. The effect is positive but smaller than temporal decay, adaptive policy learning, and the decision-agent confidence weighting, indicating that the current reference simulator is still dominated by temporal T1/T4 evidence. The module should therefore be presented as a secondary diagnostic safeguard rather than the main performance driver.

**Single-branch baselines (T1–T5 only).** Single-branch variants show that T4-only (82.3%) and T5-only (84.0%) are strong but insufficient; both trail MTIM by at least 10.4 pp in accuracy and by more than 7 pp in TSR. T3-only remains the weakest because the energy stream is deliberately easy for evasive attackers to mimic. These results support multi-factor temporal fusion, while also revealing that the current model places most discriminative power in behavior/task-history channels.

## 4.5 Scalability and Sensitivity Analysis

To evaluate how MTIM scales with network size, we varied the number of UAVs from 5 to 50 while proportionally increasing the number of Thinker Agents ( $M = \lceil N/3 \rceil$ ). Figure 4(a) shows robustness under varying malicious-node ratios at  $N = 15$ ; Table 7 reports end-to-end latency, communication overhead, and accuracy as functions of  $N$ .

**Table 7:** Scalability with swarm size under a 30% malicious-node ratio. Values are mean  $\pm$  std over 50 trials.

| UAV count $N$ | Thinker Agents $M$ | Accuracy (%)   | Eval. Time (ms) | Comm. Overhead (KB/UAV) |
|---------------|--------------------|----------------|-----------------|-------------------------|
| 5             | 2                  | $97.2 \pm 7.0$ | $36.5 \pm 3.0$  | $27.1 \pm 1.9$          |
| 15            | 5                  | $94.4 \pm 5.6$ | $58.3 \pm 4.3$  | $34.0 \pm 1.8$          |
| 30            | 10                 | $93.5 \pm 3.8$ | $95.1 \pm 4.2$  | $44.9 \pm 1.9$          |
| 50            | 17                 | $95.2 \pm 2.9$ | $152.0 \pm 3.8$ | $59.8 \pm 1.7$          |

At  $N = 50$  UAVs, the per-cycle evaluation latency grows to  $152.0 \pm 3.8$  ms, which remains well within the 10 s trust update interval in the reference setting. The dominant scaling factor is the diagnostic feature extraction and aggregation step, whose cost grows with  $M \cdot N$ . Communication overhead scales from  $34.0 \pm 1.8$  KB/UAV at  $N = 15$  to  $59.8 \pm 1.7$  KB/UAV at  $N = 50$ . Detection accuracy remains above 93% across the tested range, although the small- $N$  case has high variance because one seed-level misclassification changes the trial accuracy substantially.

Sensitivity analysis is limited to the trained policy architecture, fixed threshold, and trust-stream generator. A full grid search over policy-head loss weights, trust-decay rates, feature windows, and threshold calibration is not yet implemented. This limitation means that the reported configuration should be interpreted as one validated design point rather than a globally optimized policy.

## 4.6 Trust Factor Correlation Analysis

To verify that the five trust dimensions capture complementary rather than redundant information, we computed the pairwise Pearson correlation coefficients among  $T_i^{(1)}$  through  $T_i^{(5)}$  across all UAVs and time steps in the 50-trial simulation set. The correlation matrix (averaged over trials) is:

The strongest correlations involve the historical trust dimension T5, especially with T4 ( $0.65 \pm 0.09$ ) and T1 ( $0.64 \pm 0.08$ ), which is expected because T5 is an exponential history of recent trust evidence. Among the instantaneous dimensions T1–T4, all correlations remain at or below 0.41. Thus, T5 should be interpreted as a temporal memory channel rather than an independent sensor, while the instantaneous factors retain enough diversity to justify multi-factor fusion.

**Table 8:** Pairwise Pearson correlations among the five trust dimensions (mean  $\pm$  std over 50 trials).

|    | T1 (Behavior) | T2 (Comm.)      | T3 (Energy)     | T4 (Task)       | T5 (History)    |
|----|---------------|-----------------|-----------------|-----------------|-----------------|
| T1 | 1.00          | $0.28 \pm 0.10$ | $0.11 \pm 0.08$ | $0.41 \pm 0.10$ | $0.64 \pm 0.08$ |
| T2 | –             | 1.00            | $0.13 \pm 0.09$ | $0.27 \pm 0.10$ | $0.51 \pm 0.07$ |
| T3 | –             | –               | 1.00            | $0.11 \pm 0.09$ | $0.34 \pm 0.09$ |
| T4 | –             | –               | –               | 1.00            | $0.65 \pm 0.09$ |
| T5 | –             | –               | –               | –               | 1.00            |

## 5 Conclusions

This paper presented the Multi-Agent Trust and Incentive Mechanism (MTIM), a closed-loop framework that uses trust evaluation to regulate UAV behavior through reward and penalty decisions. The key technical contributions are: (i) a trust–incentive problem formulation that links behavior assessment, trust scoring, and task-level actions; (ii) five complementary trust dimensions with a cross-factor consistency check for detecting multi-dimensional behavioral manipulation; (iii) a multi-agent execution architecture where distributed Thinker Agents supervise local UAV subsets and an Intelligent Decision Agent fuses local estimates through confidence-weighted aggregation; and (iv) an offline-trained MARL-style trust–incentive surrogate that strengthens the mapping from trust evidence to reward/penalty actions under partial observability.

Under mixed attack patterns (constant, dormant, on-off) at a 30% malicious-node ratio, MTIM achieves 94.4% detection accuracy, 91.8% F-score, and 90.5% task success rate in the trust-layer evaluation. It significantly outperforms manual-fusion and linear risk baselines, while a strong MLP-Trust baseline obtains statistically indistinguishable detection accuracy and F-score. MTIM’s main advantage over this strong learned baseline is therefore mission-level: its TSR is 2.2 pp higher because the policy couples risk scoring to incentive-aware task allocation. Ablation studies show that temporal/adaptive evidence and incentive coupling are the main contributors, while cross-factor checking provides a secondary 1.5 pp accuracy gain and still requires stronger targeted stress tests. Scalability tests show that accuracy remains above 95% at  $N = 50$  UAVs and per-cycle latency remains within the 10 s update budget in the reference simulator. A supplementary replication package is provided for independent verification of the reported trust-layer results.

The central implication of this work is that *trust evaluation should be treated as a behavior-regulation loop, not only as an open-loop classification task*. By using trust scores to decide rewards, penalties, and task participation, MTIM links trust assessment to mission-level behavior while keeping execution-time decisions based only on observable trust evidence.

**Limitations and future work.** Several limitations warrant acknowledgment. First, MTIM mitigates only limited behavioral manipulation under fixed identities; it does not solve Sybil attacks or strong collusion in which malicious nodes dominate local observations. Second, the Intelligent Decision Agent remains a logically centralized component; a fully distributed decision protocol (e.g., based on Byzantine consensus) would strengthen the zero-trust claim. Third, the implemented MARL policy is an offline-trained trust–incentive surrogate rather than an online reinforcement-learning system with deployment-time adaptation. Fourth, the main numerical results are still produced at the trust-factor-stream layer rather than from multi-seed ns-3 packet-level experiments or hardware-in-the-loop flights. Future work will address these limitations through (i) collusion-resistant peer-evidence modeling, (ii) Sybil-resistant identity binding and admission control, (iii) decentralized Byzantine-resilient decision protocols, (iv) multi-seed ns-3 and software-defined-radio validation with public traces, and (v) safe online MARL adaptation with policy constraints.

## Author Contributions

All authors contributed to the conceptualization, methodology, analysis, and manuscript preparation. Author roles will be specified in the camera-ready version.

## Funding

This research received no external funding.

## Institutional Review Board Statement

Not applicable.

## Informed Consent Statement

Not applicable.

## Data Availability Statement

The replication package is included under `artifacts/`. It provides the code, configurations, trained models, evaluation outputs, statistical summaries, generated figures, baseline reimplementations, and packet-trace conversion materials needed to reproduce the trust-layer results reported in this manuscript. The included packet-trace validation is a single short sanity check rather than a multi-seed packet-level validation dataset.

## Conflicts of Interest

The authors declare no conflict of interest.

## References

- [1] Kundu, J.; Alam, S.; Das, J.C.; Dey, A.; Bari, S. Trust-Based Flying Ad Hoc Network: A Survey. *IEEE Access* **2024**, *12*, 99258–99281.
- [2] Noor, F.; Khan, M.A.; Al-Zahrani, A.; Ullah, I.; Al-Dhlan, K.A. A Review on Communications Perspective of Flying Ad-Hoc Networks: Key Enabling Wireless Technologies, Applications, Challenges and Open Research Topics. *Drones* **2020**, *4*, 65.
- [3] Brik, B.; Ksentini, A.; Bouaziz, M. Federated Learning for UAVs-Enabled Wireless Networks: Use Cases, Challenges, and Open Problems. *IEEE Netw.* **2020**, *34*, 144–150.
- [4] Nguyen, D.C.; Ding, M.; Pathirana, P.N.; Seneviratne, A.; Li, J.; Poor, H.V. Federated Learning for Internet of Things: A Comprehensive Survey. *IEEE Commun. Surv. Tutor.* **2021**, *23*, 1622–1658.
- [5] Qu, Y.; Dai, H.; Zhuang, Y.; Chen, J.; Dong, C.; Wu, F.; Guo, S. Decentralized Federated Learning for UAV Networks: Architecture, Challenges, and Opportunities. *IEEE Netw.* **2021**, *35*, 156–162.
- [6] Zuo, J.; Cao, R.; Qi, J.; Gao, P.; Wang, Z.; Li, J.; Zhang, L.; Lu, Y. A Hierarchical Blockchain-Based Trust Measurement Method for Drone Cluster Nodes. *Drones* **2023**, *7*, 627.
- [7] Ge, C.; Zhou, L.; Hancke, G.P.; Su, C. A Provenance-Aware Distributed Trust Model for Resilient Unmanned Aerial Vehicle Networks. *IEEE Internet Things J.* **2021**, *8*, 12481–12489.
- [8] Benfriha, S.; Labraoui, N.; Bensaid, R.; Bany Salameh, H.; Saidi, H. FUBA: A Fuzzy-Based Unmanned Aerial Vehicle Behaviour Analytics for Trust Management in Flying Ad-Hoc Networks. *IET Networks* **2024**, *13*, 208–220.
- [9] Li, H.; Li, X.; Dunkin, F.; Zhang, Z.; Lu, X. Adaptive Multi-Granularity Trust Management Scheme for UAV Visual Sensor Security under Adversarial Attacks. *Comput. Secur.* **2025**, *148*, 104108.

- [10] Akram, J.; Anaissi, A.; Rathore, R.S.; Jhaveri, R.H.; Akram, A. GALTrust: Generative Adversarial Learning-Based Framework for Trust Management in Spatial Crowdsourcing Drone Services. *IEEE Trans. Consum. Electron.* **2024**, *70*, 6196–6207.
- [11] Ceviz, O.; Sadioglu, P.; Sen, S.; Vassilakis, V.G. A Novel Federated Learning-Based IDS for Enhancing UAVs Privacy and Security. *Internet Things* **2025**, *31*, 101592.
- [12] Lu, Y.; Yang, T.; Zhao, C.; Chen, W.; Liu, Y.; Zeng, R. A Swarm Anomaly Detection Model for IoT UAVs Based on a Multi-Modal Denoising Autoencoder and Federated Learning. *Comput. Ind. Eng.* **2024**, *196*, 110454.
- [13] Gad, G.; Farrag, A.; Aboufotouh, A.; Bedda, K.; Fadlullah, Z.M.; Fouda, M.M. Joint Self-Organizing Maps and Knowledge-Distillation-Based Communication-Efficient Federated Learning for Resource-Constrained UAV-IoT Systems. *IEEE Internet Things J.* **2024**, *11*, 34270–34282.
- [14] Wang, X.; Yuan, Q.; Liu, C.; Wang, P.; Tao, J.; Li, L.; Wang, X.; Wang, Y. Who Is the Wolf in Sheep's Clothing? A Context-Aware Trust Evaluation Model for Malicious UAV Detection during Authentication. *Comput. Netw.* **2026**, *279*, 112145.
- [15] Rose, S.; Borchert, O.; Mitchell, S.; Connelly, S. Zero Trust Architecture. *NIST Spec. Publ.* **2020**, *800-207*, 1–59.
- [16] Ekechi, C.C.; Elfouly, T.; Alouani, A.; Khattab, T. A Survey on UAV Control with Multi-Agent Reinforcement Learning. *Drones* **2025**, *9*, 484.
- [17] Akram, J.; Anaissi, A.; Rathore, R.S.; Jhaveri, R.H.; Akram, A. Digital Twin-Driven Trust Management in Open RAN-Based Spatial Crowdsourcing Drone Services. *IEEE Trans. Green Commun. Netw.* **2024**, *8*, 1061–1075.
- [18] Abubaker, Z.; Javaid, N.; Almogren, A.; Akbar, M.; Zuair, M.; Ben-Othman, J. Blockchained Service Provisioning and Malicious Node Detection via Federated Learning in Scalable Internet of Sensor Things Networks. *IEEE Trans. Netw. Serv. Manag.* **2023**, *20*, 1847–1865.
- [19] Akram, J.; Umair, M.; Jhaveri, R.H.; Riaz, M.; Chi, H.R.; Malebary, S. Chained-Drones: Blockchain-Based Privacy-Preserving Framework for Secure and Intelligent Service Provisioning in Internet of Drone Things. *Comput. Electr. Eng.* **2023**, *110*, 108772.
- [20] Zheng, J.; Xu, J.; Du, H.; Niyato, D.; Kang, J.; Nie, J.; Wang, Z. Trust Management of Tiny Federated Learning in Internet of Unmanned Aerial Vehicles. *IEEE Internet Things J.* **2024**, *11*, 21046–21060.
- [21] Tong, Z.; Wang, J.; Hou, X.; Chen, J.; Jiao, Z.; Liu, J. Blockchain-Based Trustworthy and Efficient Hierarchical Federated Learning for UAV-Enabled IoT Networks. *IEEE Internet Things J.* **2024**, *11*, 15504–15522.
- [22] Xu, Q.; Li, R.; Qi, Y.; Su, Z.; Fang, D. Trust-Enhanced Game Incentive for Secure Quantum Federated Learning in UAV-Assisted Wireless Networks. *IEEE J. Sel. Areas Commun.* **2025**, *43*, 2841–2856.
- [23] Hu, Y.; Gan, Y.; Wu, H.; Wang, C.; Ma, M.; Xiong, C. GDM-DTM: A Group Decision-Making-Enabled Dynamic Trust Management Method for Malicious Node Detection in Low-Altitude UAV Networks. *Sensors* **2025**, *25*, 3982.
- [24] Wang, X.; Yuan, Q.; Wang, Y.; Teng, J.; Tao, J.; Shen, M. Zero Trust-Driven Collaborative Intrusion Detection in Internet of Things: A Continuous Trust Assessment Approach. *IEEE Internet Comput.* **2025**, *29*, 37–47.
- [25] Al Ridhawi, I.; Aloqaily, M. Zero-Trust UAV-Enabled and DT-Supported 6G Networks. *IEEE Internet Comput.* **2023**, *27*, 59–67.
- [26] Wang, L.; Wang, K.; Pan, C.; Xu, W.; Aslam, N.; Hanzo, L. Multi-Agent Deep Reinforcement Learning Based Trajectory Planning for Multi-UAV Assisted Mobile Edge Computing. *IEEE Trans. Cogn. Commun. Netw.* **2021**, *7*, 73–84.

- [27] Nie, Y.; Zhao, J.; Gao, F.; Yu, F.R. Semi-Distributed Resource Management in UAV-Aided MEC Systems: A Multi-Agent Federated Reinforcement Learning Approach. *IEEE Trans. Veh. Technol.* **2021**, *70*, 13162–13173.
- [28] Alam, M.M.; Moh, S. Joint Optimization of Trajectory Control, Task Offloading, and Resource Allocation in Air-Ground Integrated Networks. *IEEE Internet Things J.* **2024**, *11*, 24273–24288.
- [29] Yang, H.; Zhao, J.; Xiong, Z.; Lam, K.-Y.; Sun, S.; Xiao, L. Privacy-Preserving Federated Learning for UAV-Enabled Networks: Learning-Based Joint Scheduling and Resource Management. *IEEE J. Sel. Areas Commun.* **2021**, *39*, 3144–3159.
- [30] Lim, W.Y.B.; Huang, J.; Xiong, Z.; Kang, J.; Niyato, D.; Hua, X.-S.; Leung, C.; Miao, C. Towards Federated Learning in UAV-Enabled Internet of Vehicles: A Multi-Dimensional Contract-Matching Approach. *IEEE Trans. Intell. Transp. Syst.* **2021**, *22*, 5140–5154.
- [31] Wang, S.; Hosseinalipour, S.; Gorlatova, M.; Brinton, C.G.; Chiang, M. UAV-Assisted Online Machine Learning Over Multi-Tiered Networks: A Hierarchical Nested Personalized Federated Learning Approach. *IEEE Internet Things J.* **2024**, *11*, 20899–20911.
- [32] Qureshi, K.I.; Wang, L.; Xiong, X.; Lodhi, M.A. Asynchronous Federated Learning for Resource Allocation in Software-Defined Internet of UAVs. *IEEE Internet Things J.* **2024**, *11*, 15523–15538.
- [33] Jing, Y.; Qu, Y.; Dong, C.; Ren, W.; Shen, Y.; Wu, Q.; Guo, S. Exploiting UAV for Air-Ground Integrated Federated Learning: A Joint UAV Location and Resource Optimization Approach. *IEEE Trans. Green Commun. Netw.* **2023**, *7*, 1420–1433.
- [34] Fu, Z.; Liu, J.; Mao, Y.; Qu, L.; Xie, L.F.; Wang, X. Energy-Efficient UAV-Assisted Federated Learning: Trajectory Optimization, Device Scheduling, and Resource Management. *IEEE Trans. Netw. Serv. Manag.* **2025**, *22*, 974–988.
- [35] Cao, X.; Nan, G.; Guo, H.; Mu, H.; Wang, L.; Lin, Y.; Zhou, Q.; Li, J.; Qin, B.; Cui, Q.; Tao, X.; Fang, H.; Du, H.; Quek, T.Q.S. Exploring LLM-Based Multi-Agent Situation Awareness for Zero-Trust Space-Air-Ground Integrated Network. *IEEE J. Sel. Areas Commun.* **2025**, *43*, 2230–2247.
- [36] Bertsekas, D.P.; Tsitsiklis, J.N. *Neuro-Dynamic Programming*; Athena Scientific: Nashua, NH, USA, 1996.
- [37] Oliehoek, F.A.; Spaan, M.T.; Vlassis, N. Optimal and Approximate Solutions to Decentralized POMDPs. In *Proceedings of the 17th International Conference on Automated Planning and Scheduling (ICAPS)*; AAAI Press: 2008; pp. 274–281.
- [38] Riley, G.F.; Henderson, T.R. The ns-3 Network Simulator. In *Modeling and Tools for Network Simulation*; Wehrle, K.; Günes, M.; Gross, J., Eds.; Springer: Berlin, Germany, 2010; pp. 15–34.
- [39] Hausknecht, M.; Stone, P. Deep Recurrent Q-Learning for Partially Observable MDPs. In *Proceedings of the AAAI Fall Symposium on Sequential Decision Making for Intelligent Agents*; AAAI Press: 2015; pp. 29–37.
- [40] Hausknecht, M.; Stone, P. Deep Reinforcement Learning in Parameterized Action Space. In *Proceedings of the International Conference on Learning Representations (ICLR)*; 2016.
- [41] Xiong, J.; Wang, Q.; Yang, Z.; Sun, P.; Han, L.; Zheng, Y.; Fu, H.; Zhang, T.; Liu, J.; Liu, H. Parametrized Deep Q-Networks Learning: Reinforcement Learning with Discrete-Continuous Hybrid Action Space. *arXiv* **2018**, arXiv:1810.06394.